

dog_app

July 13, 2019

1 Convolutional Neural Networks

1.1 Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets:

Note: if you are using the Udacity workspace, you DO NOT need to re-download these - they can be found in the /data folder as noted in the cell below.

- Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dog_images.
- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays human_files and dog_files.

```
In [1]: import numpy as np
        from glob import glob

        # load filenames for human and dog images
        human_files = np.array(glob("/data/lfw/*/"))
        dog_files = np.array(glob("/data/dog_images/*/"))

        # print number of images in each dataset
        print('There are %d total human images.' % len(human_files))
        print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the haarcascades directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [2]: import cv2
        import matplotlib.pyplot as plt
        %matplotlib inline

        # extract pre-trained face detector
        face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

        # load color (BGR) image
        img = cv2.imread(human_files[0])
        # convert BGR image to grayscale
        gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

        # find faces in image
        faces = face_cascade.detectMultiScale(gray)

        # print number of faces detected in the image
        print('Number of faces detected:', len(faces))
```

```

# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()

```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns True if a human face is detected in an image and False otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [3]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    return len(faces) > 0
```

1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer: (You can print out your results and/or write your percentages in this cell)

Detected human face 98 out of 100 human files Detected dog face 17 out of 100 dog files:

```
In [6]: from tqdm import tqdm

human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

#-#-# Do NOT modify the code above this line. #-#-#

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.
def face_detection_perf_test(files):
    detection_cnt = 0;
    total_cnt = len(files)
    for file in files:
        detection_cnt += face_detector(file)
    return detection_cnt, total_cnt

print("Detected human face {} out of {} human files".format(face_detection_perf_test(human_files_short)[0], human_files_short.shape[0]))
print("Detected dog face {} out of {} dog files:".format(face_detection_perf_test(dog_files_short)[0], dog_files_short.shape[0]))
```

Detected human face 98 out of 100 human files

Detected dog face 17 out of 100 dog files:

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make

use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [ ]: ### (Optional)
        ### TODO: Test performance of another face detection algorithm.
        ### Feel free to use as many code cells as needed.
```

Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of 1000 categories.

```
In [7]: import torch
        import torchvision.models as models

        # define VGG16 model
        VGG16 = models.vgg16(pretrained=True)

        # check if CUDA is available
        use_cuda = torch.cuda.is_available()
        print("Is cuda available? {0}".format(use_cuda))

        # move model to GPU if CUDA is available
        if use_cuda:
            VGG16 = VGG16.cuda()
```

```
Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /root/.torch/models/vgg
100%|| 553433881/553433881 [00:07<00:00, 72138292.41it/s]
```

```
Is cuda available? True
```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as `'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg'`) as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```
In [8]: from PIL import Image
import torchvision.transforms as transforms

def VGG16_predict(img_path):
    """
    Use pre-trained VGG-16 model to obtain index corresponding to
    predicted ImageNet class for image at specified path

    Args:
        img_path: path to an image

    Returns:
        Index corresponding to VGG-16 model's prediction
    """

    ## TODO: Complete the function.
    ## Load and pre-process an image from the given img_path
    ## Return the *index* of the predicted class for that image

    image = Image.open(img_path).convert('RGB')
    # resize to (244, 244) because VGG16 accept this shape
    in_transform = transforms.Compose([
        transforms.Resize(size=(244, 244)),
        transforms.ToTensor()])

    # discard the transparent, alpha channel (that's the :3)
    image = in_transform(image)[:3,:,:].unsqueeze(0)

    if use_cuda:
        image = image.cuda()
    ret = VGG16(image)

    # return predicted class index
    return torch.max(ret,1)[1].item()

# predict dog using ImageNet class
VGG16_predict(dog_files_short[0])
```

Out [8]: 243

1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all

categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the `dog_detector` function below, which returns True if a dog is detected in an image (and False if not).

```
In [10]: def dog_detector(img_path):
         idx = VGG16_predict(img_path)
         return idx >= 151 and idx <= 268

         print("Dog Found in dog files {}".format(dog_detector(dog_files_short[0])))
         print("Dog Found in human files {}".format(dog_detector(human_files_short[0])))
```

```
Dog Found in dog files True
Dog Found in human files False
```

1.1.6 (IMPLEMENTATION) Assess the Dog Detector

Question 2: Use the code cell below to test the performance of your `dog_detector` function.

- What percentage of the images in `human_files_short` have a detected dog?
- What percentage of the images in `dog_files_short` have a detected dog?

Answer:

```
In [11]: ### TODO: Test the performance of the dog_detector function
         ### on the images in human_files_short and dog_files_short.
         def dog_detector_test(files):
             detection_cnt = 0;
             total_cnt = len(files)
             for file in files:
                 detection_cnt += dog_detector(file)
             return detection_cnt, total_cnt

         print("Detected a {} dog in {} human files".format(dog_detector_test(human_files_short)
         print("Detected a {} dog in {} dog files".format(dog_detector_test(dog_files_short)[0],
```

```
Detected a 0 dog in 100 human files
Detected a 96 dog in 100 dog files
```

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [ ]: ### (Optional)
         ### TODO: Report the performance of another pre-trained network.
         ### Feel free to use as many code cells as needed.
```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [14]: import os
         from torchvision import datasets
         import torchvision.transforms as transforms
         import torch
```



```

import numpy as np
from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True

batch_size = 20
num_workers = 0

data_dir = '/data/dog_images/'
train_dir = os.path.join(data_dir, 'train/')
valid_dir = os.path.join(data_dir, 'valid/')
test_dir = os.path.join(data_dir, 'test/')

standard_normalization = transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                                std=[0.229, 0.224, 0.225])

data_transforms = {'train': transforms.Compose([transforms.RandomResizedCrop(224),
                                                transforms.RandomHorizontalFlip(),
                                                transforms.ToTensor(),
                                                standard_normalization]),
                  'val': transforms.Compose([transforms.Resize(256),
                                                transforms.CenterCrop(224),
                                                transforms.ToTensor(),
                                                standard_normalization]),
                  'test': transforms.Compose([transforms.Resize(size=(224,224)),
                                                transforms.ToTensor(),
                                                standard_normalization])
                  }

train_data = datasets.ImageFolder(train_dir, transform=data_transforms['train'])
valid_data = datasets.ImageFolder(valid_dir, transform=data_transforms['val'])
test_data = datasets.ImageFolder(test_dir, transform=data_transforms['test'])

train_loader = torch.utils.data.DataLoader(train_data,
                                            batch_size=batch_size,
                                            num_workers=num_workers,
                                            shuffle=True)
valid_loader = torch.utils.data.DataLoader(valid_data,
                                            batch_size=batch_size,
                                            num_workers=num_workers,
                                            shuffle=False)
test_loader = torch.utils.data.DataLoader(test_data,
                                            batch_size=batch_size,
                                            num_workers=num_workers,
                                            shuffle=False)

```

```
loaders_scratch = {
    'train': train_loader,
    'valid': valid_loader,
    'test': test_loader
}
```

Question 3: Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

Answer: RandomResizedCrop & RandomHorizontalFlip are used on train data. That is used for image resizing and augmentation. Image augmentation gives randomness and that helps preventing overfitting.

Validation data is Resized to 256 and center cropped to make 224 x 224. This configuration works well with pre-trained models.

1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```
In [16]: import torch.nn as nn
import torch.nn.functional as F
import numpy as np

from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True
num_classes = 133

# define the CNN architecture
class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class
    def __init__(self):
        super(Net, self).__init__()
        ## Define layers of a CNN
        self.conv1 = nn.Conv2d(3, 32, 3, stride=2, padding=1)
        self.conv2 = nn.Conv2d(32, 64, 3, stride=2, padding=1)
        self.conv3 = nn.Conv2d(64, 128, 3, padding=1)

        # pool
        self.pool = nn.MaxPool2d(2, 2)

        # fully-connected
        self.fc1 = nn.Linear(7*7*128, 500)
        self.fc2 = nn.Linear(500, num_classes)

        # drop-out
        self.dropout = nn.Dropout(0.3)
```

```

def forward(self, x):
    ## Define forward behavior
    x = F.relu(self.conv1(x))
    x = self.pool(x)
    x = F.relu(self.conv2(x))
    x = self.pool(x)
    x = F.relu(self.conv3(x))
    x = self.pool(x)

    # flatten
    x = x.view(-1, 7*7*128)

    x = self.dropout(x)
    x = F.relu(self.fc1(x))

    x = self.dropout(x)
    x = self.fc2(x)
    return x

##-## You so NOT have to modify the code below this line. ##-##

# instantiate the CNN
model_scratch = Net()

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()

```

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

Answer: First 2 conv layers: kernel_size of 3 with stride 2, this will lead to downsize of input image by 2. Maxpooling with stride 2 is placed after 2 conv layers to help to downsize of input image by 2.

The 3rd conv layers is consist of kernel_size of 3 with stride 1, and this will not reduce input image. after final maxpooling with stride 2, the total output image size is downsized by factor of 32 and the depth will be 128. Dropout is applied to prevent overfitting.

2nd fully-connected layer is intended to produce final output size which predicts classes of breeds.

1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```
In [17]: import torch.optim as optim
```

```

### TODO: select loss function
criterion_scratch = nn.CrossEntropyLoss()

### TODO: select optimizer
optimizer_scratch = optim.SGD(model_scratch.parameters(), lr = 0.05)

```

1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath 'model_scratch.pt'.

```

In [21]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path, last_val
"""returns trained model"""
# initialize tracker for minimum validation loss
if last_validation_loss is not None:
    valid_loss_min = last_validation_loss
else:
    valid_loss_min = np.Inf

for epoch in range(1, n_epochs+1):
    # initialize variables to monitor training and validation loss
    train_loss = 0.0
    valid_loss = 0.0

    #####
    # train the model #
    #####
    model.train()
    for batch_idx, (data, target) in enumerate(loaders['train']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        ## find the loss and update the model parameters accordingly
        ## record the average training loss, using something like
        ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
        # initialize weights to zero
        optimizer.zero_grad()

        output = model(data)

        # calculate loss
        loss = criterion(output, target)

        # back prop
        loss.backward()

        # grad
        optimizer.step()

```

```

        train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

    if batch_idx % 100 == 0:
        print('Epoch %d, Batch %d loss: %.6f' %
              (epoch, batch_idx + 1, train_loss))

    #####
    # validate the model #
    #####
    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['valid']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        ## update the average validation loss
        output = model(data)
        loss = criterion(output, target)
        valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

    # print training/validation statistics
    print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
        epoch,
        train_loss,
        valid_loss
    ))

    ## TODO: save the model if validation loss has decreased
    if valid_loss < valid_loss_min:
        torch.save(model.state_dict(), save_path)
        print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
            valid_loss_min,
            valid_loss))
        valid_loss_min = valid_loss

    # return trained model
    return model

# train the model
model_scratch = train(10, loaders_scratch, model_scratch, optimizer_scratch,
                      criterion_scratch, use_cuda, 'model_scratch.pt')

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('model_scratch.pt'))

```

Epoch 1, Batch 1 loss: 4.093311

Epoch 1, Batch 101 loss: 4.079073
 Epoch 1, Batch 201 loss: 4.069899
 Epoch 1, Batch 301 loss: 4.070626
 Epoch: 1 Training Loss: 4.064924 Validation Loss: 3.897441
 Validation loss decreased (inf --> 3.897441). Saving model ...
 Epoch 2, Batch 1 loss: 3.834417
 Epoch 2, Batch 101 loss: 4.027648
 Epoch 2, Batch 201 loss: 4.042275
 Epoch 2, Batch 301 loss: 4.036431
 Epoch: 2 Training Loss: 4.035769 Validation Loss: 3.867630
 Validation loss decreased (3.897441 --> 3.867630). Saving model ...
 Epoch 3, Batch 1 loss: 3.748473
 Epoch 3, Batch 101 loss: 3.985578
 Epoch 3, Batch 201 loss: 3.981268
 Epoch 3, Batch 301 loss: 3.972224
 Epoch: 3 Training Loss: 3.977233 Validation Loss: 3.853854
 Validation loss decreased (3.867630 --> 3.853854). Saving model ...
 Epoch 4, Batch 1 loss: 3.918594
 Epoch 4, Batch 101 loss: 3.959559
 Epoch 4, Batch 201 loss: 3.913083
 Epoch 4, Batch 301 loss: 3.917247
 Epoch: 4 Training Loss: 3.915293 Validation Loss: 3.723152
 Validation loss decreased (3.853854 --> 3.723152). Saving model ...
 Epoch 5, Batch 1 loss: 3.543007
 Epoch 5, Batch 101 loss: 3.850476
 Epoch 5, Batch 201 loss: 3.851946
 Epoch 5, Batch 301 loss: 3.866573
 Epoch: 5 Training Loss: 3.869247 Validation Loss: 3.698117
 Validation loss decreased (3.723152 --> 3.698117). Saving model ...
 Epoch 6, Batch 1 loss: 4.542932
 Epoch 6, Batch 101 loss: 3.838531
 Epoch 6, Batch 201 loss: 3.838859
 Epoch 6, Batch 301 loss: 3.828118
 Epoch: 6 Training Loss: 3.830600 Validation Loss: 3.708741
 Epoch 7, Batch 1 loss: 3.532290
 Epoch 7, Batch 101 loss: 3.756562
 Epoch 7, Batch 201 loss: 3.760272
 Epoch 7, Batch 301 loss: 3.771093
 Epoch: 7 Training Loss: 3.778624 Validation Loss: 3.640538
 Validation loss decreased (3.698117 --> 3.640538). Saving model ...
 Epoch 8, Batch 1 loss: 3.900812
 Epoch 8, Batch 101 loss: 3.769214
 Epoch 8, Batch 201 loss: 3.761209
 Epoch 8, Batch 301 loss: 3.746613
 Epoch: 8 Training Loss: 3.749504 Validation Loss: 3.628509
 Validation loss decreased (3.640538 --> 3.628509). Saving model ...
 Epoch 9, Batch 1 loss: 3.781700
 Epoch 9, Batch 101 loss: 3.694990

```

Epoch 9, Batch 201 loss: 3.730481
Epoch 9, Batch 301 loss: 3.720459
Epoch: 9          Training Loss: 3.716258          Validation Loss: 3.637510
Epoch 10, Batch 1 loss: 3.506036
Epoch 10, Batch 101 loss: 3.661058
Epoch 10, Batch 201 loss: 3.668637
Epoch 10, Batch 301 loss: 3.673872
Epoch: 10         Training Loss: 3.679499          Validation Loss: 3.570432
Validation loss decreased (3.628509 --> 3.570432). Saving model ...

```

1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```

In [22]: def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['test']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
        100. * correct / total, correct, total))

    # call test function
    test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)

```

Test Loss: 3.721468

Test Accuracy: 13% (111/836)

Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dogImages/train`, `dogImages/valid`, and `dogImages/test`, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```
In [23]: ## TODO: Specify data loaders
        loaders_transfer = loaders_scratch.copy()
```

1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```
In [24]: import torchvision.models as models
        import torch.nn as nn

        ## TODO: Specify model architecture
        model_transfer = models.resnet50(pretrained=True)
        for param in model_transfer.parameters():
            param.requires_grad = False

        model_transfer.fc = nn.Linear(2048, 133, bias=True)
        fc_parameters = model_transfer.fc.parameters()

        for param in fc_parameters:
            param.requires_grad = True

        if use_cuda:
            model_transfer = model_transfer.cuda()
```

```
Downloading: "https://download.pytorch.org/models/resnet50-19c8e357.pth" to /root/.torch/models/
100%|| 102502400/102502400 [00:03<00:00, 25865893.63it/s]
```


Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer: Selected model ResNet performed best on image classification.

1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```
In [25]: criterion_transfer = nn.CrossEntropyLoss()
         optimizer_transfer = optim.SGD(model_transfer.fc.parameters(), lr=0.001)
```

1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```
In [27]: # train the model
         #model_transfer = # train(n_epochs, loaders_transfer, model_transfer, optimizer_transfer)

         # load the model that got the best validation accuracy (uncomment the line below)
         #model_transfer.load_state_dict(torch.load('model_transfer.pt'))
         def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
             """returns trained model"""
             # initialize tracker for minimum validation loss
             valid_loss_min = np.Inf

             for epoch in range(1, n_epochs+1):
                 # initialize variables to monitor training and validation loss
                 train_loss = 0.0
                 valid_loss = 0.0

                 #####
                 # train the model #
                 #####
                 model.train()
                 for batch_idx, (data, target) in enumerate(loaders['train']):
                     # move to GPU
                     if use_cuda:
                         data, target = data.cuda(), target.cuda()

                     # initialize weights to zero
                     optimizer.zero_grad()

                     output = model(data)

                     # calculate loss
                     loss = criterion(output, target)
```

```

        # back prop
        loss.backward()

        # grad
        optimizer.step()

        train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

        if batch_idx % 100 == 0:
            print('Epoch %d, Batch %d loss: %.6f' %
                  (epoch, batch_idx + 1, train_loss))

#####
# validate the model #
#####
model.eval()
for batch_idx, (data, target) in enumerate(loaders['valid']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    ## update the average validation loss
    output = model(data)
    loss = criterion(output, target)
    valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
    epoch,
    train_loss,
    valid_loss
))

## TODO: save the model if validation loss has decreased
if valid_loss < valid_loss_min:
    torch.save(model.state_dict(), save_path)
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
        valid_loss_min,
        valid_loss))
    valid_loss_min = valid_loss

# return trained model
return model

train(20, loaders_transfer, model_transfer, optimizer_transfer, criterion_transfer, use

# load the model that got the best validation accuracy

```

```

        model_transfer.load_state_dict(torch.load('model_transfer.pt'))

Epoch 1, Batch 1 loss: 2.350386
Epoch 1, Batch 101 loss: 2.284555
Epoch 1, Batch 201 loss: 2.285551
Epoch 1, Batch 301 loss: 2.277313
Epoch: 1          Training Loss: 2.273011          Validation Loss: 1.686224
Validation loss decreased (inf --> 1.686224).  Saving model ...
Epoch 2, Batch 1 loss: 2.185141
Epoch 2, Batch 101 loss: 2.244076
Epoch 2, Batch 201 loss: 2.235444
Epoch 2, Batch 301 loss: 2.224876
Epoch: 2          Training Loss: 2.221145          Validation Loss: 1.618313
Validation loss decreased (1.686224 --> 1.618313).  Saving model ...
Epoch 3, Batch 1 loss: 2.097821
Epoch 3, Batch 101 loss: 2.153676
Epoch 3, Batch 201 loss: 2.175415
Epoch 3, Batch 301 loss: 2.153899
Epoch: 3          Training Loss: 2.154629          Validation Loss: 1.536086
Validation loss decreased (1.618313 --> 1.536086).  Saving model ...
Epoch 4, Batch 1 loss: 2.014405
Epoch 4, Batch 101 loss: 2.117808
Epoch 4, Batch 201 loss: 2.123878
Epoch 4, Batch 301 loss: 2.122859
Epoch: 4          Training Loss: 2.119215          Validation Loss: 1.497218
Validation loss decreased (1.536086 --> 1.497218).  Saving model ...
Epoch 5, Batch 1 loss: 2.112144
Epoch 5, Batch 101 loss: 2.071465
Epoch 5, Batch 201 loss: 2.069074
Epoch 5, Batch 301 loss: 2.048761
Epoch: 5          Training Loss: 2.050658          Validation Loss: 1.473037
Validation loss decreased (1.497218 --> 1.473037).  Saving model ...
Epoch 6, Batch 1 loss: 1.885552
Epoch 6, Batch 101 loss: 1.994177
Epoch 6, Batch 201 loss: 2.011742
Epoch 6, Batch 301 loss: 1.999248
Epoch: 6          Training Loss: 2.004949          Validation Loss: 1.401448
Validation loss decreased (1.473037 --> 1.401448).  Saving model ...
Epoch 7, Batch 1 loss: 1.413045
Epoch 7, Batch 101 loss: 1.980999
Epoch 7, Batch 201 loss: 1.986720
Epoch 7, Batch 301 loss: 1.978199
Epoch: 7          Training Loss: 1.975926          Validation Loss: 1.357703
Validation loss decreased (1.401448 --> 1.357703).  Saving model ...
Epoch 8, Batch 1 loss: 1.943936
Epoch 8, Batch 101 loss: 1.937140
Epoch 8, Batch 201 loss: 1.944394
Epoch 8, Batch 301 loss: 1.936126

```

Epoch: 8 Training Loss: 1.931879 Validation Loss: 1.309746
 Validation loss decreased (1.357703 --> 1.309746). Saving model ...
 Epoch 9, Batch 1 loss: 1.951793
 Epoch 9, Batch 101 loss: 1.901610
 Epoch 9, Batch 201 loss: 1.898619
 Epoch 9, Batch 301 loss: 1.900769
 Epoch: 9 Training Loss: 1.897734 Validation Loss: 1.279616
 Validation loss decreased (1.309746 --> 1.279616). Saving model ...
 Epoch 10, Batch 1 loss: 1.804905
 Epoch 10, Batch 101 loss: 1.854092
 Epoch 10, Batch 201 loss: 1.899708
 Epoch 10, Batch 301 loss: 1.887789
 Epoch: 10 Training Loss: 1.885328 Validation Loss: 1.246197
 Validation loss decreased (1.279616 --> 1.246197). Saving model ...
 Epoch 11, Batch 1 loss: 1.856759
 Epoch 11, Batch 101 loss: 1.817694
 Epoch 11, Batch 201 loss: 1.835331
 Epoch 11, Batch 301 loss: 1.833365
 Epoch: 11 Training Loss: 1.830395 Validation Loss: 1.197896
 Validation loss decreased (1.246197 --> 1.197896). Saving model ...
 Epoch 12, Batch 1 loss: 1.514580
 Epoch 12, Batch 101 loss: 1.796287
 Epoch 12, Batch 201 loss: 1.797184
 Epoch 12, Batch 301 loss: 1.803827
 Epoch: 12 Training Loss: 1.801712 Validation Loss: 1.186543
 Validation loss decreased (1.197896 --> 1.186543). Saving model ...
 Epoch 13, Batch 1 loss: 2.438747
 Epoch 13, Batch 101 loss: 1.752473
 Epoch 13, Batch 201 loss: 1.758562
 Epoch 13, Batch 301 loss: 1.756177
 Epoch: 13 Training Loss: 1.759254 Validation Loss: 1.159904
 Validation loss decreased (1.186543 --> 1.159904). Saving model ...
 Epoch 14, Batch 1 loss: 1.710798
 Epoch 14, Batch 101 loss: 1.776765
 Epoch 14, Batch 201 loss: 1.759389
 Epoch 14, Batch 301 loss: 1.742555
 Epoch: 14 Training Loss: 1.742235 Validation Loss: 1.150779
 Validation loss decreased (1.159904 --> 1.150779). Saving model ...
 Epoch 15, Batch 1 loss: 2.161588
 Epoch 15, Batch 101 loss: 1.711720
 Epoch 15, Batch 201 loss: 1.749598
 Epoch 15, Batch 301 loss: 1.721810
 Epoch: 15 Training Loss: 1.721262 Validation Loss: 1.080289
 Validation loss decreased (1.150779 --> 1.080289). Saving model ...
 Epoch 16, Batch 1 loss: 1.682941
 Epoch 16, Batch 101 loss: 1.672051
 Epoch 16, Batch 201 loss: 1.672442
 Epoch 16, Batch 301 loss: 1.669082

```

Epoch: 16          Training Loss: 1.672202          Validation Loss: 1.113229
Epoch 17, Batch 1 loss: 1.475960
Epoch 17, Batch 101 loss: 1.634007
Epoch 17, Batch 201 loss: 1.642253
Epoch 17, Batch 301 loss: 1.645759
Epoch: 17          Training Loss: 1.643476          Validation Loss: 1.068408
Validation loss decreased (1.080289 --> 1.068408). Saving model ...
Epoch 18, Batch 1 loss: 1.626438
Epoch 18, Batch 101 loss: 1.586108
Epoch 18, Batch 201 loss: 1.622340
Epoch 18, Batch 301 loss: 1.643316
Epoch: 18          Training Loss: 1.645823          Validation Loss: 1.053460
Validation loss decreased (1.068408 --> 1.053460). Saving model ...
Epoch 19, Batch 1 loss: 2.015040
Epoch 19, Batch 101 loss: 1.625617
Epoch 19, Batch 201 loss: 1.611728
Epoch 19, Batch 301 loss: 1.609136
Epoch: 19          Training Loss: 1.606944          Validation Loss: 1.005818
Validation loss decreased (1.053460 --> 1.005818). Saving model ...
Epoch 20, Batch 1 loss: 1.554231
Epoch 20, Batch 101 loss: 1.588878
Epoch 20, Batch 201 loss: 1.598705
Epoch 20, Batch 301 loss: 1.597323
Epoch: 20          Training Loss: 1.596734          Validation Loss: 1.008483

```

1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [28]: test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

```
Test Loss: 1.091834
```

```
Test Accuracy: 79% (667/836)
```

1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
In [32]: ### TODO: Write a function that takes a path to an image as input
        ### and returns the dog breed that is predicted by the model.
```

```
from PIL import Image
```

```

import torchvision.transforms as transforms

# list of class names by index, i.e. a name can be accessed like class_names[0]
#class_names = [item[4:].replace("_", " ") for item in data_transfer['train'].classes]
class_names = [item[4:].replace("_", " ") for item in loaders_transfer['train'].dataset]

def load_input_image(img_path):
    image = Image.open(img_path).convert('RGB')
    prediction_transform = transforms.Compose([transforms.Resize(size=(224, 224)),
                                              transforms.ToTensor(),
                                              standard_normalization])

    # discard the transparent, alpha channel (that's the :3) and add the batch dimension
    image = prediction_transform(image)[:3,:,:].unsqueeze(0)
    return image

def predict_breed_transfer(model, class_names, img_path):
    # load the image and return the predicted breed
    img = load_input_image(img_path)
    model = model.cpu()
    model.eval()
    idx = torch.argmax(model(img))
    return class_names[idx]

for img_file in os.listdir('./images'):
    img_path = os.path.join('./images', img_file)
    predition = predict_breed_transfer(model_transfer, class_names, img_path)
    print("{0} predition breed: {1}".format(img_path, predition))

./images/Welsh_springer_spaniel_08203.jpg predition breed: Welsh springer spaniel
./images/Labrador_retriever_06457.jpg predition breed: Labrador retriever
./images/Labrador_retriever_06455.jpg predition breed: Chesapeake bay retriever
./images/sample_human_output.png predition breed: Brussels griffon
./images/Brittany_02625.jpg predition breed: Brittany
./images/American_water_spaniel_00648.jpg predition breed: Curly-coated retriever
./images/Labrador_retriever_06449.jpg predition breed: Flat-coated retriever
./images/Curly-coated_retriever_03896.jpg predition breed: Curly-coated retriever
./images/sample_cnn.png predition breed: Brussels griffon
./images/sample_dog_output.png predition breed: Great dane

```

Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.



Sample Human Output

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.


Some sample output for our algorithm is provided below, but feel free to design your own user experience!

1.1.18 (IMPLEMENTATION) Write your Algorithm


```
In [34]: ### TODO: Write your algorithm.
         ### Feel free to use as many code cells as needed.

def run_app(img_path):
    ## handle cases for a human face, dog, and neither
    img = Image.open(img_path)
    plt.imshow(img)
    plt.show()
    if dog_detector(img_path) is True:
        prediction = predict_breed_transfer(model_transfer, class_names, img_path)
        print("Dogs Detected!\nIt looks like a {}".format(prediction))
    elif face_detector(img_path) > 0:
        prediction = predict_breed_transfer(model_transfer, class_names, img_path)
        print("Hello, human!\nIf you were a dog..You may look like a {}".format(prediction))
    else:
        print("Error! Can't detect anything..")


for img_file in os.listdir('./images'):
    img_path = os.path.join('./images', img_file)
    run_app(img_path)
```



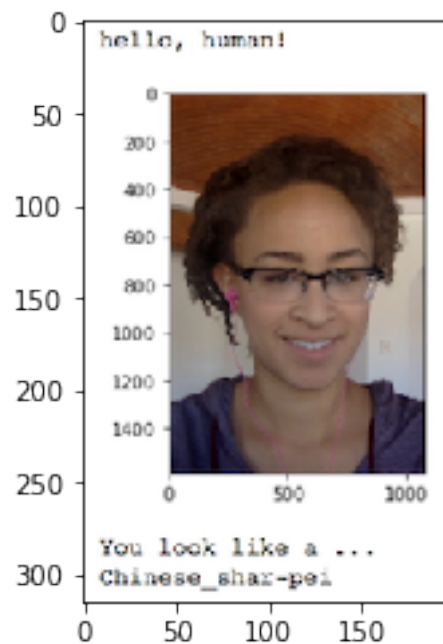
Dogs Detected!
It looks like a Welsh springer spaniel



Dogs Detected!
It looks like a Labrador retriever




Dogs Detected!
It looks like a Chesapeake bay retriever



Hello, human!

If you were a dog..You may look like a Brussels griffon



Dogs Detected!


It looks like a Brittany




Dogs Detected!
It looks like a Curly-coated retriever



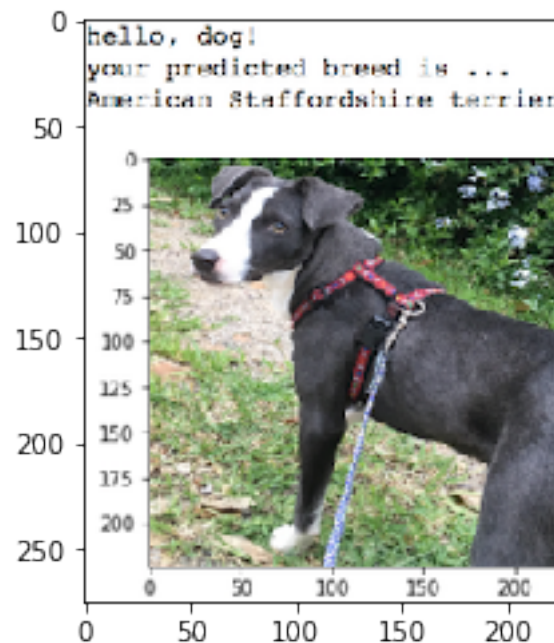
Dogs Detected!
It looks like a Flat-coated retriever



Dogs Detected!
It looks like a Curly-coated retriever



Error! Can't detect anything..



Dogs Detected!
It looks like a Great dane

Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer: (Three possible points for improvement) Areas of improvement 1. With bigger training data we can improve accuracy. Current model only applied RandomResizedCrop & RandomHorizontalFlip but with more augmentations like Vertical Flips will improve accuracy. With current mode may miss images with flipped images.


2. Test with couple of models. Current model worked well for single dog in image. Need to test it with multiple dogs in same image. 3. Can be optimized and tested with more hyper parameters.

```
In [44]: ## TODO: Execute your algorithm from Step 6 on
        ## at least 6 images on your computer.
        ## Feel free to use as many code cells as needed.
        human_files = ['./test_images/tesla.jpeg', './test_images/Ramanujan.jpeg', './test_images/']
        dog_files = ['./test_images/Pug.jpeg', './test_images/Labrador.jpg', './test_images/Pom']

        import os
        print(os.getcwd())

        ## suggested code, below
        for file in np.hstack((human_files[:3], dog_files[:3])):
            run_app(file)


/home/workspace/dog_project
```



Hello, human!
 If you were a dog..You may look like a American foxhound

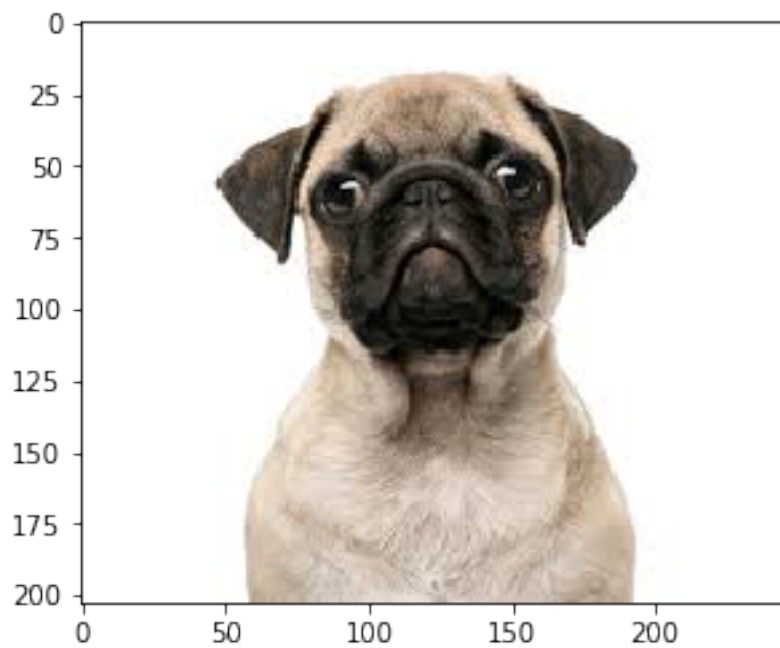


Hello, human!
If you were a dog..You may look like a Greyhound



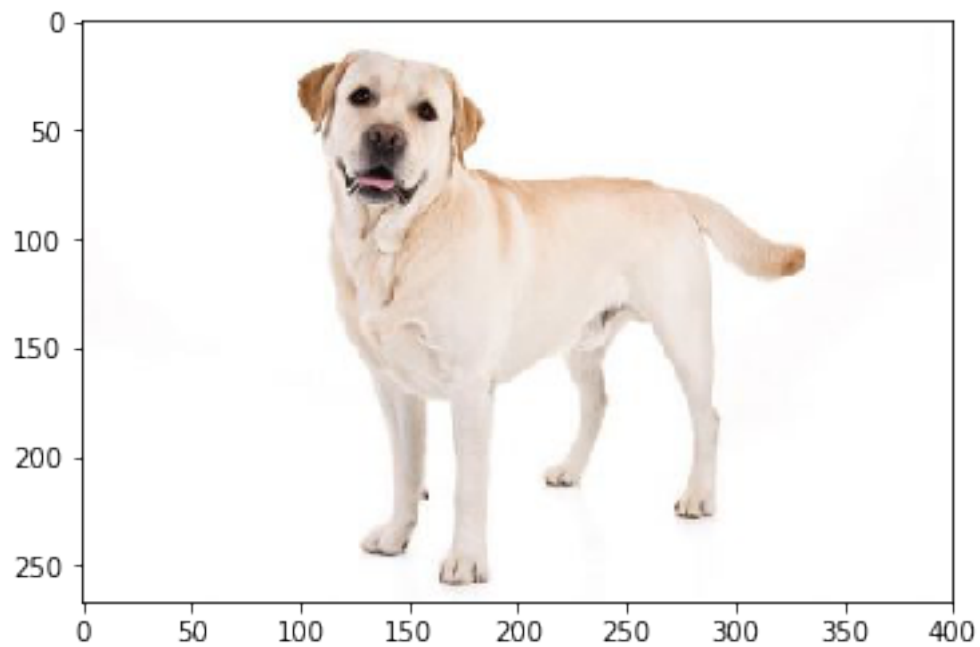
Hello, human!

If you were a dog..You may look like a Irish wolfhound



Dogs Detected!

It looks like a Bullmastiff



Dogs Detected!

It looks like a Labrador retriever



Dogs Detected!

It looks like a American eskimo dog